



Chapter 7.

Transaction Management and Concurrency Control

Transaction Support
Properties of Transaction Support
Concurrency Control
Concurrency Control Techniques
 Deadlock
 Database Recovery
 Recovery Facilities
 Recovery Techniques

Copyright © Aini Yunita Abdul Rahman 2017



Introduction

- Reliability and consistency must be maintained in the presence of failures of both hardware and software components and when multiple users are accessing the database.
- All functions are dependent.
- Both concurrency control and recovery services are required to protect database from data inconsistencies and data loss.
- Many DBMS allow simultaneously operation on the database.
- If these operations are not controlled, the accesses will interfere with one another and the database become inconsistent.

Copyright © Aini Yunita Abdul Rahman 2017



■ To overcome this problem, DBMS implements a concurrency control protocol that prevent databases accesses from interfering with one another.

■ Database recovery is the process of restoring the database to a correct state following a failure.

■ The failure caused by the system crash due to hardware or software errors in the application or media failure such as head crash.

■ Whatever cause of failure, the DBMS must be able to recover from failure and restore the database to consistent state.

Copyright © Aini Yunita Abdul Rahman 2017



Transaction Support

Transaction

Action or series of actions, carried out by user or application, which reads or updates contents of database.

- Logical unit of work on the database, may be involve entire program or single SQL statement or any operation.
- Application program is series of transactions with non-database processing in between.
- Transforms database from one consistent state to another, although consistency may be violated during transaction.

Copyright © Aini Yunita Abdul Rahman 2017



Example Transaction

```

read(staffNo = x, salary)
salary = salary * 1.1
write(staffNo = x, new_salary)

(a)                                     (b)

      delete(staffNo = x)
      for all PropertyForRent records, pno
      begin
          read(propertyNo = pno, staffNo)
          if (staffNo = x) then
              begin
                  staffNo = newStaffNo
                  write(propertyNo = pno, staffNo)
              end
      end
  
```

Copyright © Aini Yunita Abdul Rahman 2017

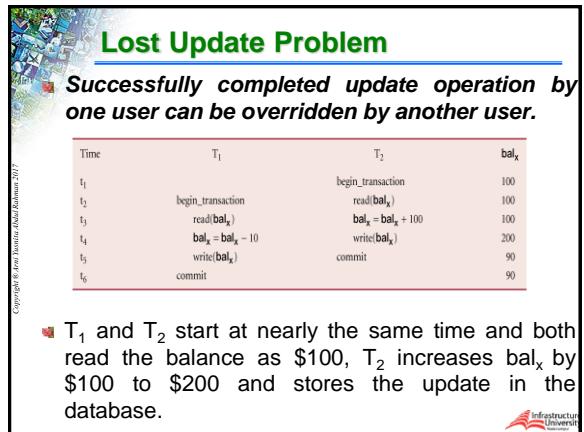
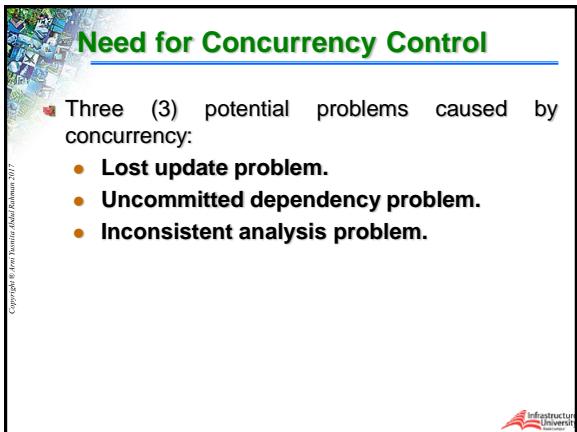
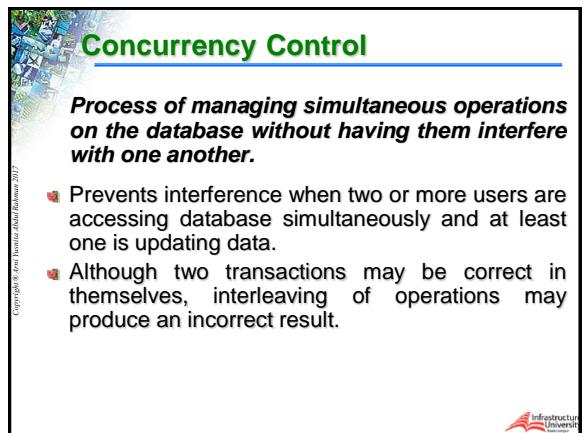
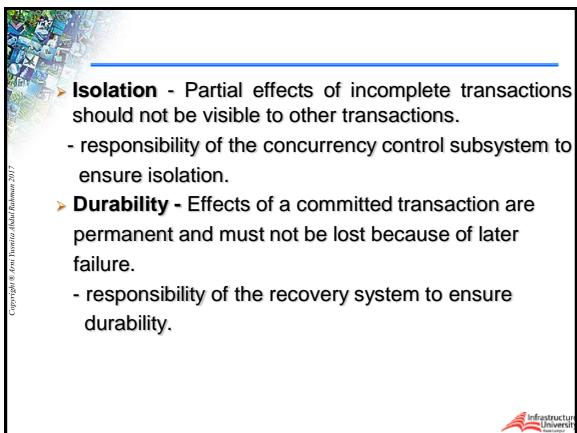
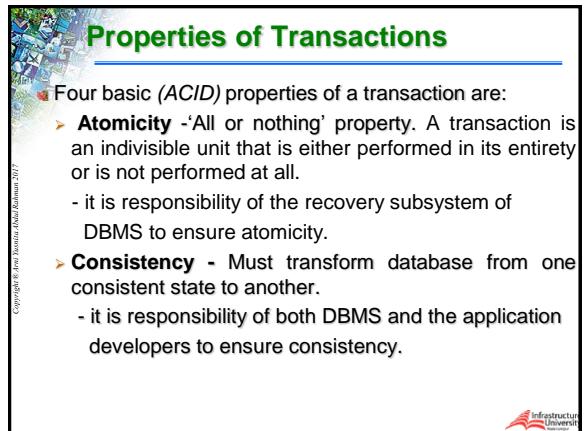
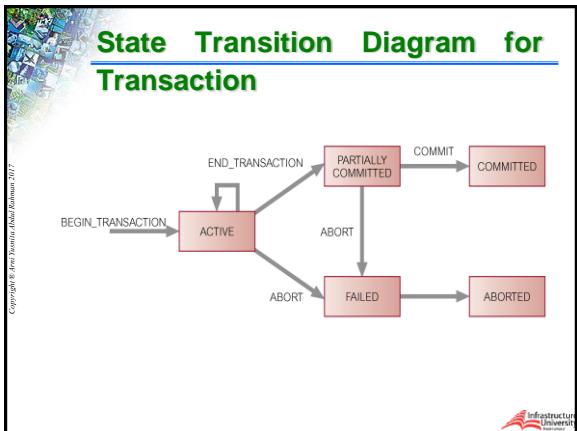


Can have one of two outcomes:

- **Success** - transaction *commits* and database reaches a new consistent state.
- **Failure** - transaction *aborts* and database must be restored to consistent state before it started.
 - Such a transaction is *rolled back* or *undone*.
 - Committed transaction cannot be aborted.
 - If committed transaction is a mistake, then have to perform compensating transaction to reverse its effect.
 - Aborted transaction that is rolled back can be restarted later.

Copyright © Aini Yunita Abdul Rahman 2017





Copyright © from [Ivan Iordanoff](#) 2017

Meanwhile transaction T_1 decrements its copy of bal_x by \$10 to \$90 and stores this value in the database, overwriting the previous update and thereby losing the \$100 previously added to the balance.

- Loss of T_2 's update avoided by preventing T_1 from reading bal_x until after update.

Copyright © from [Ivan Iordanoff](#) 2017

Uncommitted Dependency Problem

- Occurs when one transaction can see intermediate results of another transaction before it has committed.
- T_4 updates bal_x to \$200 but it aborts, so bal_x should be back at original value of \$100.
- T_3 has read new value of bal_x (\$200) and uses value as basis of \$10 reduction, giving a new balance of \$190, instead of \$90.

Copyright © from [Ivan Iordanoff](#) 2017

Time	T_3	T_4	bal_x
t_1		begin_transaction	100
t_2		read(bal_x)	100
t_3		$\text{bal}_x = \text{bal}_x + 100$	100
t_4	begin_transaction	write(bal_x)	200
t_5	read(bal_x)	:	200
t_6	$\text{bal}_x = \text{bal}_x - 10$	rollback	100
t_7		write(bal_x)	190
t_8	commit		190

Problem avoided by preventing T_3 from reading bal_x until after T_4 commits or aborts.

Copyright © from [Ivan Iordanoff](#) 2017

Inconsistent Analysis Problem

- Occurs when transaction reads several values but second transaction updates some of them during execution of first.
- Sometimes referred to as dirty read or unrepeatable read.
- T_6 is totaling balances of account x (\$100), account y (\$50) and account z (\$25).
- Meantime, T_5 has transferred \$10 from bal_x to bal_z , so T_6 now has wrong result (\$10 too high).

Copyright © from [Ivan Iordanoff](#) 2017

Time	T_5	T_6	bal_x	bal_y	bal_z	sum
t_1		begin_transaction	100	50	25	175
t_2	begin_transaction	sum = 0	100	50	25	0
t_3	read(bal_x)	read(bal_x)	100	50	25	0
t_4	$\text{bal}_x = \text{bal}_x - 10$	sum = sum + bal_x	100	50	25	100
t_5	write(bal_x)	read(bal_y)	90	50	25	100
t_6	read(bal_x)	sum = sum + bal_x	90	50	25	150
t_7	$\text{bal}_x = \text{bal}_x + 10$	sum = sum + bal_y	90	50	25	150
t_8	write(bal_z)	sum = sum + bal_z	90	50	35	150
t_9	commit	read(bal_x)	90	50	35	150
t_{10}		sum = sum + bal_z	90	50	35	185
t_{11}	commit		90	50	35	185

Problem avoided by preventing T_6 from reading bal_x and bal_z until after T_5 completed updates.

Copyright © from [Ivan Iordanoff](#) 2017

Concurrency Control Techniques

- Two basic concurrency control techniques:
 - Locking,
 - Timestamping.
- Both are conservative approaches: delay transactions in case they conflict with other transactions.

Locking

Transaction uses locks to deny access to other transactions and to prevent incorrect updates.

- Generally, a transaction must claim a *shared* (read) or *exclusive* (write) lock on a data item before read or write.
- Lock prevents another transaction from modifying item or even reading it, in the case of a write lock.

Copyright © from [Institutul Politehnica Bucuresti](#) 2017

Locking - Basic Rules

- If transaction has shared lock on item, can read but not update item.
- If transaction has exclusive lock on item, can both read and update item.
- Reads cannot conflict, so more than one transaction can hold shared locks simultaneously on same item.
- Exclusive lock gives transaction exclusive access to that item.
- Some systems allow transaction to upgrade read lock to an exclusive lock or downgrade exclusive lock to a shared lock.

Copyright © from [Institutul Politehnica Bucuresti](#) 2017

Deadlock

An impasse that may result when two (or more) transactions are each waiting for locks held by the others to be released.

Time	T ₁₇	T ₁₈
t ₁	begin_transaction	
t ₂	write_lock(bal_x)	
t ₃	read(bal_x)	
t ₄	$bal_x = bal_x - 10$	
t ₅	write(bal_x)	
t ₆	write_lock(bal_y)	
t ₇	WAIT	
t ₈	WAIT	
t ₉	WAIT	
t ₁₀	⋮	
t ₁₁	⋮	begin_transaction
		write_lock(bal_y)
		read(bal_y)
		$bal_y = bal_y + 100$
		write(bal_y)
		write_lock(bal_x)
		WAIT
		WAIT
		WAIT
		⋮

Copyright © from [Institutul Politehnica Bucuresti](#) 2017

- Only one way to break deadlock: abort one or more of the transactions.
- Deadlock should be transparent to user, so DBMS should restart transaction(s).
- Three general techniques for handling deadlock:
 - Timeouts.
 - Deadlock prevention.
 - Deadlock detection and recovery.

Copyright © from [Institutul Politehnica Bucuresti](#) 2017

Timeouts

Transaction that requests lock will only wait for a system-defined period of time.

- If lock has not been granted within this period, lock request times out.
- In this case, DBMS assumes transaction may be deadlocked, even though it may not be and it aborts and automatically restarts the transaction.

Copyright © from [Institutul Politehnica Bucuresti](#) 2017

Deadlock Prevention

- DBMS looks ahead to see if transaction would cause deadlock and never allows deadlock to occur.
- Could order transactions using transaction timestamps:
 - Wait-Die - only an older transaction can wait for younger one, otherwise transaction is aborted (*dies*) and restarted with same timestamp.
 - Wound-Wait - only a younger transaction can wait for an older one. If older transaction requests lock held by younger one, younger one is aborted (*wounded*).

Copyright © from [Institutul Politehnica Bucuresti](#) 2017

Deadlock Detection and Recovery

DBMS allows deadlock to occur but recognizes it and breaks it.

- Usually handled by construction of Wait-For-Graph (WFG) showing transaction dependencies:
 - Create a node for each transaction.
 - Create edge $T_i \rightarrow T_j$, if T_i waiting to lock item locked by T_j .
- Deadlock exists if and only if WFG contains cycle.

Wait-For-Graph (WFG)

Copyright © from Author/Infra Roman 2017

Timestamping

- Transactions ordered globally so that older transactions, transactions with smaller timestamps, get priority in the event of conflict.
- Conflict is resolved by rolling back and restarting transaction.
- No locks so no deadlock.

Timestamp

A unique identifier created by DBMS that indicates relative starting time of a transaction.

- Can be generated by using system clock at time transaction started or by incrementing a logical counter every time a new transaction starts.

Copyright © from Author/Infra Roman 2017

Database Recovery

Process of restoring database to a correct state in the event of a failure.

Need for Recovery Control

- Two types of storage: volatile (main memory) and nonvolatile.
- Volatile storage does not survive system crashes.
- Stable storage represents information that has been replicated in several nonvolatile storage media with independent failure modes.

Copyright © from Author/Infra Roman 2017

Types of Failures

- System crashes** due to hardware or software errors, resulting in loss of main memory.
- Media failures** such as head crashes or unreadable media, resulting in loss of parts of secondary storage.
- Application software errors** such as logical errors in the program that is accessing the database which cause one or more transactions to fail.
- Natural physical disasters** such as fires, floods, earthquakes or power failures.
- Carelessness or unintentional destruction** of data or facilities by operators or users.
- Sabotage** or intentional corruption or destruction of data, hardware or software facilities.

Copyright © from Author/Infra Roman 2017

Recovery Facilities

DBMS should provide following facilities to assist with recovery:

- Backup mechanism**, which makes periodic backup copies of database.
- Logging facilities**, which keep track of current state of transactions and database changes.
- Checkpoint facility**, which enables updates to database in progress to be made permanent.
- Recovery manager**, which allows DBMS to restore database to consistent state following a failure.

Copyright © from Author/Infra Roman 2017

Recovery Techniques

If database has been damaged:

- Need to restore last backup copy of database and reapply updates of committed transactions using log file.

If database is only inconsistent:

- Need to undo changes that caused inconsistency. May also need to redo some transactions to ensure updates reach secondary storage.
- No need to backup but can restore database using before-and-after-images in the log file.

Copyright © from Author/Infra Roman 2017